Edge Video Analytics

Abstract

Edge video analytics has emerged as a pivotal paradigm in modern computing, enabling real-time intelligent analysis of visual data at or near the point of capture rather than in distant cloud data centres. The convergence of Edge AI, deep learning, and the Internet of Things (IoT) has unlocked transformative capabilities across smart cities, industrial safety, surveillance, and autonomous systems. This write-up synthesises findings from five recent peer-reviewed publications spanning 2022–2025 to provide a structured overview of the technological foundations, key application domains, pressing challenges, and promising future directions in edge video analytics. Together, these works underline a consistent narrative: as camera deployments proliferate and latency requirements tighten, edge-based intelligence is no longer optional but essential.

1. Introduction

The global proliferation of cameras — from smart city surveillance nodes and industrial floor monitors to autonomous vehicle sensors and wearable devices — has generated an unprecedented torrent of video data. Traditional architectures that route raw video streams to centralised cloud servers for analysis are increasingly untenable: they impose high bandwidth costs, introduce unacceptable latency, and raise significant privacy concerns. Edge computing addresses these limitations by positioning computational resources close to data sources, at routers, gateways, on-camera chips, or nearby micro-servers.

Edge video analytics refers to the deployment of AI-driven video analysis pipelines on such edge hardware. By running inference — and even selective training — locally, edge systems can deliver sub-second responses, operate under limited or intermittent connectivity, and process sensitive footage without transmitting identifiable data to remote servers. Gartner has projected that 50% of all AI inference will occur at the edge by 2025, and research confirms that edge video analytics is entering a period of rapid maturity with transformative potential [2].

This review draws on five recent peer-reviewed publications to trace the landscape of edge video analytics from its technological underpinnings through to its real-world applications, ongoing challenges, and future research directions.

2. Technological Foundations

2.1 Edge Computing Architecture

At its core, edge computing extends cloud resources — compute, storage, and networking — to the immediate vicinity of the devices that generate and consume data [5]. This creates a three-tier hierarchy: IoT/terminal devices (cameras, sensors), edge nodes (on-device chips, micro-servers, base stations), and the cloud. Each tier handles tasks suited to its computational capacity. Edge nodes serve as intermediaries that pre-process video streams, execute lightweight AI inference, and selectively forward aggregated results or anomalous segments upstream.

Edge analytics further refines this model by embedding analytical algorithms directly within edge devices or their proximate nodes [5]. The approach enables descriptive, diagnostic, and even predictive analytics at the network periphery, dramatically reducing the volume of data transmitted to cloud infrastructure and enabling decisions within milliseconds of an event occurring.

2.2 Deep Learning at the Edge

Modern edge video analytics is underpinned by deep neural networks (DNNs), particularly convolutional neural networks (CNNs) for object detection and classification, and recurrent or transformer-based architectures for

temporal understanding. However, state-of-the-art models such as those winning the ImageNet and PASCAL VOC challenges carry a computational cost that is rarely available on resource-constrained edge hardware [4].

To bridge this gap, several model optimisation strategies have been developed:

- Model quantisation: reducing the numerical precision of weights (e.g., from 32-bit floating point to 8-bit integers) to lower memory footprint and inference latency.
- Pruning: removing redundant neural network connections to produce lighter models with minimal accuracy loss.
- Knowledge distillation: training a compact 'student' model to mimic a larger 'teacher' model.
- Model specialisation: fine-tuning a general lightweight model on scene-specific data to exceed the accuracy of the base model within a narrower scope [4].

The COVA (Contextually Optimised Video Analytics) framework exemplifies model specialisation: it automatically generates per-camera training datasets from video streams, fine-tunes object detection models using TensorFlow, and deploys the resulting specialised models to edge cameras [4]. This pipeline automates what would otherwise be a labour-intensive customisation process, enabling scalable deployment across large networks of cameras.

2.3 Hardware Accelerators

The practical viability of edge video analytics depends heavily on purpose-built silicon. Platforms such as the NVIDIA Jetson family (including the Jetson Xavier NX used in industrial safety experiments [3]), Google Coral TPUs, and dedicated NPUs (Neural Processing Units) provide orders-of-magnitude improvements in inference throughput and energy efficiency compared to general-purpose CPUs. Recent processor generations leverage sub-5nm manufacturing processes to balance computational performance against power budgets, which is critical for battery-powered or thermally constrained deployments [1].

2.4 Edge–Cloud Collaboration

A recurrent theme across the reviewed literature is that edge video analytics does not eliminate the cloud but rather reconfigures the relationship. Edge nodes handle latency-sensitive tasks — real-time detection, alerting, and privacy filtering — while the cloud supports computationally heavy functions such as long-horizon model training, cross-site data aggregation, and system-wide coordination [2]. Hybrid architectures partition computation dynamically: systems like Planner automatically decompose stream analytics into edge and cloud segments based on resource availability and network cost models [5].

3. Application Domains

3.1 Smart City Management

Smart cities represent the most extensively studied context for edge video analytics [1][2]. Applications span traffic flow optimisation, pedestrian monitoring, parking management, crowd density estimation, and license plate recognition. Edge nodes co-located with existing CCTV infrastructure allow municipalities to retrofit legacy camera networks with AI capabilities without replacing hardware or redesigning networks.

A systematic review of edge-AI-enabled video analytics in smart cities [2] documents how deep learning pipelines running on edge servers can classify vehicles, count pedestrians, detect incidents, and estimate crowd densities at the city scale, all in real time. The review highlights data-efficient and opportunistic analytics techniques that maintain acceptable performance even when labelled training data are scarce — a common constraint in municipal deployments.

3.2 Intelligent Surveillance and Security

Surveillance is one of the oldest motivations for video analytics and has been reinvigorated by edge AI. Traditional closed-circuit systems require human operators to watch dozens of feeds simultaneously — an inherently error-prone approach. Edge AI enables autonomous detection of anomalous events such as intrusion, violence, or abandoned objects, triggering alerts with millisecond latency [3].

Edge AI in real-time surveillance systems [3] describes architectures in which AI inference runs directly on smart cameras or nearby edge servers, enabling responses before network round-trips to the cloud would even complete. Such systems reduce false-positive rates by incorporating contextual scene understanding, distinguishing, for instance, between a person legitimately working late and an unauthorised intruder.

3.3 Industrial Worker Safety

The integration of Industry 4.0 technologies — IoT devices, robotics, and automated production lines — has improved manufacturing efficiency but introduced new worker safety challenges. Traditional surveillance systems lacked real-time hazard detection and proactive intervention capabilities [3][6].

The ICTACT Journal study [6] proposed an **Edge-AI powered video analytics framework** designed to operate directly on local manufacturing site devices. Using an NVIDIA Jetson Xavier NX for edge inference, the system performed real-time PPE (Personal Protective Equipment) compliance checks, equipment malfunction detection, and unsafe-behaviour classification through a multi-stage decision-making framework. Experimental results demonstrated high-accuracy hazard scoring with alert response times well below the thresholds tolerable for human safety.

This use case illustrates a broader trend: edge video analytics is becoming a primary pillar of occupational health and safety systems, not merely a supplementary monitoring tool.

3.4 Autonomous Vehicles and Transportation

Autonomous vehicles represent perhaps the most demanding application context: they require real-time inference from multiple high-resolution camera feeds, operating under strict safety constraints with no tolerance for network latency or intermittent connectivity. Edge processing is therefore not optional but architecturally mandatory. Edge video analytics frameworks support tasks ranging from lane detection and pedestrian recognition to traffic sign classification and collision prediction [2][5].

3.5 Healthcare and Environmental Monitoring

Healthcare applications include patient monitoring in intensive care units, fall detection for elderly residents, and surgical workflow analysis. Environmental applications encompass wildlife monitoring, flood detection, and agricultural inspection. In both domains, edge deployment addresses connectivity limitations (remote hospitals, rural environments) and privacy concerns (patient data must not leave clinical premises) [1][5].

4. Challenges

4.1 Resource Constraints

Edge devices operate under tight budgets of compute, memory, and energy. The reviewed literature consistently identifies this as the primary barrier to deploying state-of-the-art deep learning models at the edge [4][5]. While model optimisation techniques mitigate the problem, they typically involve accuracy trade-offs that must be carefully managed for safety-critical applications. Determining the optimal accuracy-efficiency frontier for a given deployment context remains an open research problem.

4.2 Model Scalability and Maintenance

A large-scale smart city deployment may encompass thousands of cameras, each operating in a distinct visual context. Manually curating training data and fine-tuning models for each camera is infeasible. The COVA framework [4] addresses this through automation, but generalising such pipelines to diverse hardware platforms, camera configurations, and scene types at production scale remains challenging. Model drift — whereby real-world conditions evolve away from training distributions — further requires continuous monitoring and periodic retraining.

4.3 Reliability and Fault Tolerance

Edge nodes are physically distributed and may fail due to hardware faults, power outages, or network disruptions. Robust architectures must ensure that failure of individual nodes does not compromise the broader analytics system [5]. Failover mechanisms, redundant edge nodes, and graceful degradation to reduced-functionality modes are active areas of investigation. Edge analytics frameworks must also handle the asynchronous, heterogeneous data streams produced by diverse sensor networks.

4.4 Privacy and Data Governance

While edge processing inherently reduces privacy risk by avoiding raw video transmission to the cloud, local edge nodes themselves must be secured. Sensitive footage — in workplaces, hospitals, or residential areas — may be exposed to physical tampering, side-channel attacks, or unauthorised local access. Regulatory frameworks such as GDPR impose obligations regarding data minimisation, purpose limitation, and the right to erasure, creating compliance complexity for deployments that aggregate facial or behavioural data even locally [1][3].

Privacy-by-design approaches — including on-device anonymisation, differential privacy, and federated learning — are gaining traction as responses to these concerns, but their integration into real-time edge video analytics pipelines at acceptable computational cost is still maturing.

4.5 Interoperability and Standardisation

Edge video analytics deployments frequently combine hardware from multiple vendors, operating systems, and AI frameworks. The absence of standardised interfaces for model deployment, telemetry, and management creates integration burdens and vendor lock-in risks [5]. Containerisation technologies (Docker, Kubernetes for the edge) partially address this through abstraction, but comprehensive interoperability standards for edge AI in video analytics remain immature.

4.6 Security

Edge devices deployed in public or industrial environments are exposed to adversarial threats, including model inversion attacks, adversarial examples designed to fool object detectors, and physical tampering with cameras or edge hardware [1]. Unlike cloud infrastructure, edge nodes cannot always rely on enterprise-grade physical and network security, requiring lightweight on-device security mechanisms that do not exhaust computational budgets.

5. Future Directions

5.1 Federated Learning

Federated learning enables multiple edge nodes to collaboratively train a shared model without centralising raw data. Nodes compute local model updates from their private video streams and transmit only gradient information to a central aggregator. This approach harmonises the competing demands of model generalisation (which benefits from diverse data) and data privacy (which requires data to remain local) [1][3]. Integrating federated learning into production edge video analytics systems — particularly under communication constraints and with heterogeneous device capabilities — is a rich research frontier.

5.2 Neuromorphic and Energy-Efficient Hardware

Neuromorphic processors, which model computation on the sparse, event-driven processing of biological neurons, offer the potential for dramatic reductions in inference energy consumption [1]. Unlike conventional GPUs that process full video frames at fixed frame rates, event-based cameras and neuromorphic chips respond only to changes in the scene, reducing redundant computation. As these technologies mature, they may unlock edge video analytics in battery-powered and energy-harvesting contexts currently beyond reach.

5.3 Generative AI and Foundation Models

Large vision-language foundation models (e.g., CLIP, GPT-4V) possess broad visual understanding that could be fine-tuned for edge video analytics tasks with minimal task-specific data. Techniques such as parameter-efficient fine-tuning and model distillation are being explored to compress these models for edge deployment. Generative AI also opens new possibilities for multi-stage, natural-language-driven querying of video archives — enabling operators to search footage semantically rather than through rigid rule-based filters [1].

5.4 Explainable and Trustworthy AI

As edge video analytics systems take consequential decisions — triggering emergency stops, alerting law enforcement, or flagging workers for non-compliance — explainability becomes critical. Organisations, regulators, and affected individuals increasingly require not only that a decision was made but also why. Research into lightweight explainability methods suitable for edge deployment (e.g., attention visualisation, saliency maps) and into formal verification of AI safety properties in embedded contexts is accelerating [1][6].

5.5 5G and Next-Generation Connectivity

The deployment of 5G networks substantially reduces the latency and increases the bandwidth of wireless links between edge nodes and cloud infrastructure, enabling more fluid dynamic partitioning of computation. Multi-access Edge Computing (MEC) standards, which embed compute resources within 5G base stations, promise to extend the reach of edge video analytics to mobile and vehicular contexts. Research is exploring how to exploit these improved connectivity characteristics to support adaptive quality streaming and collaborative multi-camera analytics [2][5].

5.6 Automated Model Lifecycle Management

Beyond the initial specialisation problem addressed by COVA [4], there is a need for automated systems that monitor model performance in deployment, detect drift, trigger retraining, and manage the rollout of updated models across heterogeneous edge fleets — all with minimal human intervention. This encompasses version control for models, A/B testing at the edge, and intelligent scheduling of resource-intensive retraining tasks.

6. Synthesis and Discussion

The five reviewed publications collectively paint a consistent picture. Edge video analytics is not a single technology but an ecosystem of hardware accelerators, optimised AI models, communication protocols, and application frameworks. Its value proposition rests on four pillars: low latency (enabling real-time response), reduced bandwidth (processing locally rather than streaming raw video), enhanced privacy (keeping sensitive data on-premises), and improved reliability (operating under intermittent connectivity).

The progression from general-purpose edge computing [5] through smart city applications [1][2] to domain-specific deployments in surveillance [3] and industrial safety [6] reflects a maturing field. Early research established feasibility; recent work addresses the harder problems of scalability, automation (COVA [4]), and safety-critical deployment (ICTACT [6]).

A persistent tension runs through the literature: the trade-off between model accuracy and computational efficiency. This trade-off is not merely technical — it is contextual. A smart city traffic counter can tolerate occasional misclassifications; a worker safety system cannot. Future research must develop principled, application-aware methodologies for navigating this trade-off, ideally with automated tooling to calibrate the optimal operating point for each deployment context.

Privacy and security, while acknowledged as challenges in all reviewed works, remain under-addressed in practical deployment frameworks. The field would benefit from greater integration of privacy-engineering principles into core edge video analytics architectures, rather than treating privacy as an afterthought.

7. Conclusion

Edge video analytics represents one of the most consequential intersections of AI, systems engineering, and societal infrastructure in the current decade. By bringing intelligence to the camera, to the factory floor, and to the city corner, it enables applications that were either technically infeasible or economically unviable under purely cloud-centric architectures.

The reviewed literature establishes a strong foundation: robust theoretical frameworks for edge analytics [5], empirical demonstrations of real-time performance in smart cities [1][2] and surveillance [3], automated model specialisation pipelines [4], and safety-focused industrial deployments [6]. Yet significant challenges remain — in resource efficiency, model lifecycle management, privacy, security, and standardisation — ensuring that edge video analytics will remain an active and impactful research area for years to come.

As hardware accelerators grow more capable, connectivity improves with 5G and beyond, and AI models become more efficient through techniques such as federated learning and knowledge distillation, the boundary between edge and cloud intelligence will continue to blur. The ultimate vision — pervasive, real-time, privacy-preserving visual intelligence deployed at every node of our physical world — is coming progressively within reach.

References

1. Warade, E. et al. (2025). A Critical Review of Edge AI for Real-Time Video Analytics in Smart Cities: Technological Foundations, Applications, Challenges, and Future Directions. International Journal of Research Publication and Reviews, 6(5). https://ijrpr.com/uploads/V6ISSUE5/IJRPR44559.pdf
2. Nayak, S., Patgiri, R., Waikhom, L., & Ahmed, A. (2024). A review on edge analytics: Issues, challenges, opportunities, promises, future directions, and applications. Digital Communications and Networks (Elsevier), 10(3), 783–804. https://doi.org/10.1016/j.dcan.2022.10.016
3. IJRITCC. (2023). Edge AI for Real-Time Video Analytics in Surveillance Systems. International Journal of Research in Innovative Technology and Creative Communication. https://ijritcc.org/index.php/ijritcc/article/view/8947
4. Rivas, D., Guim, F., Polo, J., Silva, P. M., Berral, J. L., & Carrera, D. (2022). Towards automatic model specialization for edge video analytics. Future Generation Computer Systems (Elsevier), 134, 399–413. https://doi.org/10.1016/j.future.2022.03.039
5. Badidi, E. et al. (2023). Opportunities, Applications, and Challenges of Edge-AI Enabled Video Analytics in Smart Cities: A Systematic Review. IEEE Access. https://doi.org/10.1109/ACCESS.2023.3300658
6. Venkataramana Sagar, G. & Ambigaipriya, S. (2025). Edge-AI Video Analytics for Industrial Worker Safety and Hazard Detection in Smart Manufacturing Environments through Decision-Making Frameworks. ICTACT Journal on Image and Video Processing, 16(2), 3752–3757. https://doi.org/10.21917/ijivp.2025.0531